

A Fixed-Income Analytics Engine: Mathematical Foundations, Numerical Verification, Yield-Curve Modelling and Interest-Rate Risk

Samuel Remedios

A computational finance dissertation

Educational computational study.
No institutional affiliation or submission is asserted.

September 2026

Study status and declaration note

This dissertation studies an existing educational computational framework. Its inputs are invented offline observations; no institutional affiliation, submission or empirical market validity is asserted. The independent internal verification record remains available as `paper/verification_report.tex`.

Abstract

Fixed-rate bond valuation requires consistency between payment timing, discounting and the interest-rate variable being changed. This dissertation examines calculations for regular bonds, smooth term-structure representation and deterministic curve risk. Nominal coupon-compounded valuation is distinguished from discounting with continuously compounded zero rates. Analytical differentiation establishes duration and convexity; a six-parameter Nelson–Siegel–Svensson function represents illustrative maturity-rate observations. Independent cash-flow calculations, high-precision exponential references, finite differences and a partition-of-unity proof assess the implementation. All 234 deterministic tests pass. The illustrative 2030 bond has a clean price of 97.876838 per 100 face value and modified duration of 4.231518 years. The curve-fit root mean squared error is 0.782762 basis points. For the longer risk bond, parallel sensitivity is 0.1697264000 currency units per one-basis-point bump; summed key-rate sensitivity differs by approximately minus 2.93×10^{-8} currency units, consistent with a cubic correction. The principal limitation is that fitted illustrative rates are assigned a zero-rate interpretation without inference from instrument prices. The evidence supports mathematical consistency under explicit conventions; empirical market validity remains untested.

Contents

Study status and declaration note	i
Abstract	ii
1 Introduction	1
2 Fixed-income foundations	2
2.1 Yield to maturity versus zero/spot rates versus par yields	2
2.2 Literature and methodological position	3
3 Fixed-rate bond analytics	4
3.1 Cash flows and valuation	4
3.2 Duration, convexity and DV01	4
4 Nelson–Siegel–Svensson modelling	6
4.1 Calibration and interpretation	6
5 Interest-rate and curve risk	8
5.1 Key-rate decomposition	8
6 Computational methodology	10
7 Mathematical and numerical verification	11
8 Illustrative numerical results	12

9 Discussion and limitations	17
10 Conclusion	18
References	19
A Detailed mathematical derivations	21
A.1 The par-bond identity	21
A.2 Price derivatives and duration	22
A.3 Central differences and DV01	22
A.4 NSS limits and stable loadings	23
A.5 Key-rate basis and partition of unity	24
A.6 First-order and finite-bump reconciliation	24
A.7 Independent numerical constructions	25
B Reproducibility	27
B.1 Execution record	27
B.2 Environment and numerical reproduction	27
B.3 Calibration settings	28
B.4 Figure and numerical provenance	28
B.5 LaTeX compilation	28
B.6 Bibliographic and editorial verification	29

1 Introduction

Bonds finance borrowing through promised payments. Their valuation depends on when those payments occur and how future amounts are discounted [1]. A single discount yield cannot generally represent rates that vary by maturity; a yield curve records rates against maturity under a stated convention [1]. Price levels alone do not establish rate sensitivity. Nelson–Siegel–Svensson (NSS) modelling represents such a curve with a six-parameter function of maturity, while deterministic interest-rate risk measures price changes under specified rate shocks rather than assigning probabilities to them.

This dissertation examines fixed-rate valuation, Nelson–Siegel–Svensson modelling and deterministic interest-rate risk through four questions:

1. How do explicit conventions determine bond prices and yield sensitivities?
2. How can a smooth parametric curve—a curve specified by a finite set of parameters—be fitted reliably?
3. How can parallel exposures to the same rate move at every maturity, shaped exposures to maturity-dependent moves and key-rate exposures localised around selected maturities be represented consistently?
4. What independent evidence supports computational correctness?

The scope comprises regular option-free bonds, whose contracts contain no rights to alter payments, and illustrative offline rates. Credit events, forecasting, trading and backtesting are excluded. The contribution is independent verification of established mathematics, not new financial theory.

Chapters 2 to 5 develop the financial argument; chapters 6 to 8 present computation and evidence. Chapters 9 and 10 delimit the conclusions; appendices A and B supply derivations and reproduction instructions.

2 Fixed-income foundations

Money available today can earn interest, so payment dates affect economic value. Present value expresses a future payment in today's units; its discount factor is today's value of one future unit. Discrete compounding credits interest periodically; a nominal annual rate is divided by the compounding frequency. Continuous compounding uses exponential accumulation [1].

A fixed-rate bond is a debt contract paying fixed interest coupons and returning principal at maturity. Face, or par, value is the principal reference amount; redemption repays it. The annual coupon rate determines annual interest, divided into coupon amounts according to payment frequency [1]. Settlement is the valuation and transfer date; maturity is the final payment date. Accrued interest allocates the current coupon to elapsed time. Dirty price includes this accrual; clean price subtracts it [1].

2.1 Yield to maturity versus zero/spot rates versus par yields

Yield to maturity (YTM) is the single rate equating discounted promised payments to dirty price under specified compounding. A zero-coupon bond pays only at redemption: its spot, or zero-coupon, rate applies to that single horizon [1]. For maturity $T > 0$ years and continuous annual decimal zero rate $r(T)$, the dimensionless discount factor is exactly

$$DF(T) = \exp[-r(T)T]. \quad (2.1)$$

A par yield is the annual coupon rate making a specified bond worth face value under given discount factors, conventionally at a coupon-date settlement with no accrual. It depends on the whole payment schedule, whereas a spot rate discounts one maturity. The term structure describes rates across maturities, or tenors; a yield curve represents it under a stated rate convention [1].

Treasury constant-maturity (CMT) yields are read from the U.S. Treasury's par yield curve and differ from zero rates [2]. Fitting them does not by itself bootstrap zeros; bootstrapping infers discount factors from instrument prices and cash flows [1, 3]. Here the invented observations have no intrinsic compounding convention; risk deliberately interprets fitted rates as continuous zeros, an educational assumption.

All rates are decimal: $0.05 = 5\%$. One basis point (bp) equals 0.01 percentage point, or 0.0001 decimal [1].

2.2 Literature and methodological position

The literature separates curve representation from estimation using instruments. Nelson–Siegel provides a compact curve family; Svensson adds flexibility through a second curvature component [4, 5]. This flexibility is useful for representing shapes, but it does not establish that six parameters are necessary for these ten illustrative observations. Institutional approaches also differ: Anderson and Sleath use a spline-based method, while Gürkaynak, Sack and Wright estimate a discount function from Treasury instruments [6, 7]. Such studies address instrument selection and estimation choices absent here. The present study assesses an existing NSS implementation; it neither compares competing curve families nor establishes empirical superiority. Its research contribution is therefore reproducible verification under specified conventions.

3 Fixed-rate bond analytics

3.1 Cash flows and valuation

Assume fixed nonnegative coupons, positive redemption and annual or semiannual payments. Dates are maturity minus multiples of 12 or 6 months, clipping nonexistent days. Settlement-date payments are excluded.

Let $F > 0$ be face value in currency, $c \geq 0$ the annual decimal coupon rate, $m \in \{1, 2\}$ payments per year, and y_{YTM} nominal annual decimal YTM compounded m times. The dimensionless accumulation factor is $B = 1 + y_{\text{YTM}}/m > 0$. For settlement d_s , preceding/on-settlement coupon d_- , next coupon d_+ , and calendar-day count days, define accrued-period fraction α , currency coupon C , and accrued interest AI :

$$C = \frac{Fc}{m}, \quad \alpha = \frac{\text{days}(d_s - d_-)}{\text{days}(d_+ - d_-)}, \quad AI = C\alpha. \quad (3.1)$$

Index the N future payments by $i = 1, \dots, N$, with dates d_i . Their coupon-period distances $q_i = i - \alpha$ are fractional between coupon dates. Cash flow CF_i equals C , plus F at redemption; PV_i denotes its currency present value. Dirty and clean currency prices $P_{\text{dirty}}, P_{\text{clean}}$ follow exactly from fractional compounding:

$$P_{\text{dirty}} = \sum_{i=1}^N \underbrace{CF_i B^{-q_i}}_{PV_i}, \quad P_{\text{clean}} = P_{\text{dirty}} - AI. \quad (3.2)$$

Par means price equals face; premium and discount mean above and below it. At coupon-date settlement, $c = y_{\text{YTM}}$ implies $P_{\text{dirty}} = P_{\text{clean}} = F$ (section A.1); between coupons exact par is not guaranteed.

3.2 Duration, convexity and DV01

Duration measures local rate sensitivity; Macaulay duration D_{Mac} is also weighted payment time [8]. Define coupon-based years $T_i = q_i/m$ and abbreviate dirty price by P ; primes denote YTM derivatives. Differentiation yields modified duration D_{Mod} :

$$D_{\text{Mac}} = \frac{\sum_i T_i PV_i}{P}, \quad P' = - \sum_i CF_i \frac{q_i}{m} B^{-q_i-1}, \quad (3.3)$$

$$D_{\text{Mod}} = -\frac{P'}{P} = \frac{D_{\text{Mac}}}{B}. \quad (3.4)$$

Both durations are conventionally expressed in years [1]. For a small decimal yield change δy , relative price change is approximately $-D_{\text{Mod}}\delta y$.

For larger yield changes, the approximation also uses price curvature. Normalised convexity $\mathcal{C}$, conventionally expressed in years squared [1], follows from

$$P'' = \sum_i CF_i \frac{q_i(q_i + 1)}{m^2} B^{-q_i - 2}, \quad \mathcal{C} = P''/P. \quad (3.5)$$

Positive redemption makes $\mathcal{C} > 0$. Taylor's local polynomial approximation gives

$$\frac{P(y_{\text{YTM}} + \delta y) - P(y_{\text{YTM}})}{P(y_{\text{YTM}})} \approx -D_{\text{Mod}}\delta y + \frac{1}{2}\mathcal{C}(\delta y)^2. \quad (3.6)$$

DV01 measures the currency price response to one basis point [1], here using down-minus-up central repricing with decimal bump $h = 10^{-4}$:

$$\text{DV01}_{\text{YTM}} = \frac{P(y_{\text{YTM}} - h) - P(y_{\text{YTM}} + h)}{2} = hPD_{\text{Mod}} + O(h^3). \quad (3.7)$$

Here $O(h^3)$ denotes a remainder whose absolute value is bounded locally by a constant times h^3 . Central repricing retains odd-order nonlinear terms beyond duration; the convexity term cancels between the symmetric prices. Both bumped yields must remain valid. Sections A.2 and A.3 provide the derivations.

4 Nelson–Siegel–Svensson modelling

The Nelson–Siegel model combines level, slope-related and curvature contributions. Svensson’s extension adds another curvature contribution, giving the Nelson–Siegel–Svensson (NSS) model [4, 5]. The parameter vector is $\theta = (\beta_0, \beta_1, \beta_2, \beta_3, \tau_1, \tau_2)$. The β coefficients are decimal rates and the positive decay scales τ_1, τ_2 are measured in years. A factor loading is a function of maturity multiplying a coefficient. For either decay scale τ , dimensionless loadings define the generic decimal rate r_{NSS} :

$$L_1(T; \tau) = \frac{1 - e^{-T/\tau}}{T/\tau}, \quad L_2(T; \tau) = L_1(T; \tau) - e^{-T/\tau}, \quad (4.1)$$

$$r_{\text{NSS}}(T; \theta) = \beta_0 + \beta_1 L_1(T; \tau_1) + \beta_2 L_2(T; \tau_1) + \beta_3 L_2(T; \tau_2). \quad (4.2)$$

Asymptotic limits, the values approached at boundaries, follow from the exponential expansions in section A.4:

$$r_{\text{NSS}}(0^+) = \beta_0 + \beta_1, \quad r_{\text{NSS}}(\infty) = \beta_0. \quad (4.3)$$

Thus β_0 sets the long limit, β_1 the short offset, and β_2, β_3 overlapping hump/trough contributions; decay scales locate the individual loadings, not necessarily turning points of their sum. These are descriptive roles, not structurally unique economic factors.

For dimensionless $x = T/\tau$, subtracting nearly equal rounded numbers in $1 - e^{-x}$ causes cancellation [9]. The implementation uses `expm1`, accurately evaluating $e^{-x} - 1$, and separate fifth-degree loading expansions below $x = 10^{-3}$.

4.1 Calibration and interpretation

Calibration selects parameters fitting observations. For n observations (u_ℓ, z_ℓ) , indexed by ℓ , maturities are years and rates decimal. Residuals e_ℓ are observed minus fitted:

$$e_\ell = z_\ell - r_{\text{NSS}}(u_\ell; \theta), \quad S(\theta) = \sum_{\ell=1}^n e_\ell^2. \quad (4.4)$$

Least squares minimises S , measured in squared decimal rates. At the selected fit $\hat{\theta}$, root mean squared error is

$$\text{RMSE} = \sqrt{S(\hat{\theta})/n}, \quad \text{RMSE}_{\text{bp}} = 10^4 \text{RMSE}. \quad (4.5)$$

Nonlinear optimisation searches for parameters minimising S ; local minima minimise it within a neighbourhood, without ensuring global optimality. Six deterministic starts reduce starting-point

dependence (chapter 6). Identification is the ability of the observations and model restrictions to distinguish parameter values; it is weak when materially different values fit almost equally well. A small fitting error does not imply unique identification [10]: a constant curve leaves decay scales undetermined. If $\tau_1 = \tau_2$, only the sum $\beta_2 + \beta_3$ enters their shared loading. Near equality can therefore produce unstable coefficients even when fitted rates change little. The generic rate r_{NSS} acquires zero-rate meaning only through the assumption in section 2.1.

5 Interest-rate and curve risk

The curve-risk calculation discounts each contractual payment separately. Actual/365 times t_i count settlement-to-payment calendar days, including leap days, divided by 365. Zero-discounted present values V_i and dirty currency price $P_{\text{zero}}[r]$ are

$$t_i = \frac{\text{days}(d_i - d_s)}{365}, \quad V_i = CF_i e^{-r(t_i)t_i}, \quad P_{\text{zero}}[r] = \sum_i V_i. \quad (5.1)$$

This uses continuous compounding and the assigned zero interpretation (section 2.1), separately from coupon-based T_i and YTM pricing.

An additive decimal shock $s(T)$ gives shocked rate $r_s(T) = r(T) + s(T)$. Full repricing recalculates discounted payments; scenario profit and loss (P&L) is currency change $\Delta P_s = P_{\text{zero}}[r+s] - P_{\text{zero}}[r]$. For constant parallel amplitude δr ,

$$\frac{dP_{\text{zero}}[r + \delta r]}{d\delta r} = - \sum_i t_i V_i e^{-\delta r t_i} < 0. \quad (5.2)$$

Positive redemption ensures higher rates lower price.

Steepening increases the long-minus-short maturity rate difference; flattening decreases it. The specified steepener assigns $(-10, 0, +10)$ bp at $(2, 10, 30)$ years; the flattener reverses these shifts. Piecewise-linear interpolation joins neighbouring anchors by straight lines, holding endpoint shifts constant outside them. These definitions are illustrative, not universal.

5.1 Key-rate decomposition

Exposure varies by payment horizon [11]. Choose J ordered key maturities k_j in years, indexed by $j = 1, \dots, J$. Dimensionless basis weight $b_j(T)$ equals one at its own key and zero at others, interpolating linearly with constant endpoint tails. Neighbouring weights sum to one, a partition of unity:

$$b_j(k_j) = 1, \quad \sum_{j=1}^J b_j(T) = 1, \quad 0 \leq b_j(T) \leq 1. \quad (5.3)$$

Key-rate DV01 measures currency sensitivity to hb_j ; parallel DV01 uses constant h :

$$\text{KRDV01}_j = \frac{P_{\text{zero}}[r - hb_j] - P_{\text{zero}}[r + hb_j]}{2}, \quad (5.4)$$

$$\text{DV01}_{\parallel} = \frac{P_{\text{zero}}[r - h] - P_{\text{zero}}[r + h]}{2}. \quad (5.5)$$

Partition makes first-order exposures additive. Independent finite bumps leave a cubic reconciliation difference $\Delta_{\text{KR}} = \sum_j \text{KRDV01}_j - \text{DV01}_{\parallel}$, derived in section A.6. Exact finite-bump equality is therefore unnecessary.

6 Computational methodology

The implementation uses Python 3.12+. Immutable dataclasses preserve contractual inputs, while `bonds.py`, `curves.py` and `risk.py` separate nominal valuation, curve fitting and zero discounting. This makes the distinct conventions visible; callers must still supply the correct rate interpretation. NumPy supplies arrays and linear initialisation; SciPy supplies nonlinear `least_squares` [12, 13]. The trust-region reflective algorithm restricts successive local steps and respects parameter bounds [14–16]. Six fixed decay starts precede joint fitting; the successful candidate with the smallest residual sum of squares is selected (section B.3). Conditioning describes how sensitively the exact solution changes under small input perturbations. Gilli and colleagues identify local optima and poor conditioning as separate difficulties and study a broader search heuristic [10]. These difficulties remain possible with the engine’s finite set of local starts.

Deterministic tests use fixed inputs to check identities, monotonicity, synthetic fits, invariants and boundaries. Finite differences approximate derivatives from displaced prices; central differences use symmetric displacements. Numerical tolerances specify acceptable discrepancies; truncation error comes from omitted approximation terms. Offline execution and recorded library versions support reproducibility (appendix B).

7 Mathematical and numerical verification

The independent internal report, `paper/verification_report.tex`, supplies separately constructed cash flows, 90-digit exponential references and analytical identities. Proofs establish results under assumptions; tests provide finite numerical evidence.

Par and clean/dirty identities agree to rounding precision. Independent central repricing agrees with the engine; YTM DV01 exceeds the duration approximation by 1.93×10^{-9} currency units. Tables 7.1 and 7.2 show derivative agreement and initial second-order convergence, followed by cancellation as steps shrink.

Table 7.1: Independent normalised derivative checks. Errors are absolute differences and differences divided by the absolute analytical reference.

Quantity; step	Analytical	Numerical	Abs. error	Rel. error
$D_{\text{Mod}}; 10^{-5}$	4.2315176920	4.2315176940	1.973×10^{-9}	4.662×10^{-10}
$\mathcal{C}; 10^{-4}$	21.1104448496	21.1104452368	3.872×10^{-7}	1.834×10^{-8}

Duration and its absolute error are in years; convexity and its absolute error in years squared. Steps are decimal YTM. Source: independent verification report, numerical duration and convexity checks.

Table 7.2: Finite-difference convergence and rounding deterioration. Both normalised central formulas have $O(h^2)$ truncation error.

Decimal step h	Absolute duration error (years)	Absolute convexity error (years squared)
10^{-2}	1.951×10^{-3}	5.912×10^{-3}
10^{-3}	1.951×10^{-5}	5.911×10^{-5}
10^{-4}	1.951×10^{-7}	3.872×10^{-7}
10^{-5}	1.973×10^{-9}	1.810×10^{-6}
10^{-6}	5.499×10^{-10}	1.247×10^{-4}
10^{-7}	2.131×10^{-9}	1.832×10^{-1}
10^{-8}	2.805×10^{-8}	4.492×10^{-1}

Source: independent verification report; explicit ten-payment reference bond, rather than engine-generated reference prices.

NSS limits agree with equation (4.3); isolated loading comparisons reach maximum relative error 4.47×10^{-13} near the numerical switch. Forty independently generated noiseless observations give RMSE 3.29×10^{-18} decimal and maximum error 1.39×10^{-17} over 1,001 comparison maturities.

Independent zero-curve pricing gives 103.23571061465249 currency units, matching the engine. Shock signs satisfy equation (5.2). The basis proof covers all maturities; 100,008 sampled points have maximum computed partition error zero. Halving the reconciliation bump reduces its discrepancy approximately eightfold (section A.6). Fresh execution passes all 234 tests.

8 Illustrative numerical results

Fresh engine calculations use the specifications in table 8.1 and illustrative term-structure observations from `examples/illustrative_curve.csv`. Table 8.2 identifies a discount bond because its clean price is below face value; its coupon is also below YTM. Modified duration implies approximately 0.423% dirty-price loss for a 10 bp rise; positive convexity moderates that loss.

Table 8.1: Illustrative bond specifications. Currency amounts refer to face value 100.

Input	YTM bond	Curve-risk bond
Settlement	15 April 2025	15 April 2025
Maturity	15 January 2030	15 January 2055
Face value	100	100
Annual coupon; frequency	4%; semiannual	4%; semiannual
Nominal annual YTM	4.5%	Not used
Discount convention	Fractional coupon periods	Continuous zeros; Actual/365

Source: author's calculations using the illustrative offline data supplied with the analytics engine.

Table 8.2: Bond analytics for the 2030 maturity under nominal YTM compounding.

Measure	Unit	Value
Accrued interest	Currency	0.994475
Dirty price	Currency	98.871313
Clean price	Currency	97.876838
Macaulay duration	Years	4.326727
Modified duration	Years	4.231518
Convexity	Years squared	21.110445
YTM DV01	Currency for 1 bp	0.041838

Source: author's calculations using the illustrative offline data supplied with the analytics engine.

Tables 8.3 and 8.4 and figure 8.1 show a sub-basis-point RMSE and a trough followed by rising rates, describing only these invented observations. Parameter precision does not establish identification.

Table 8.3: NSS fitted parameters. Coefficients are decimal rates; decay scales are years.

Coefficient	Estimate	Decay scale	Estimate
β_0	0.03959735	τ_1	1.628798
β_1	0.00406525	τ_2	16.193767
β_2	-0.01904196		
β_3	0.00913973		

Source: author's calculations using the illustrative offline data supplied with the analytics engine.

Table 8.4: Calibration quality for ten illustrative observations. Residuals are observed minus fitted.

Statistic	Value
RMSE (decimal)	0.00007827624
RMSE (bp)	0.782762
Maximum absolute residual (bp)	1.520383
Successful starts	6/6
Selected termination	Status 2: <code>ftol</code>
Selected function evaluations	17

Source: author's calculations using the illustrative offline data supplied with the analytics engine.




Figure 8.1: Illustrative maturity-rate observations and fitted NSS curve over the observed range. The percentage axis displays generic rates and has no implied zero-rate convention.

Source: author's calculations using the illustrative offline data supplied with the analytics engine.

Higher discount rates lower value (table 8.5); positive curve-price curvature makes the -25 bp gain larger than the $+25$ bp loss. Exposure concentrates at the 30-year key (figure 8.2 and table 8.6). The small negative reconciliation has the predicted cubic sign.

Table 8.5: Curve scenario prices and instantaneous P&L for the 2055 bond. Fitted rates are interpreted as continuous zeros.

Scenario	Base	Shocked	P&L
Parallel $+25$ bp	98.699959	94.580653	-4.119307
Parallel -25 bp	98.699959	103.072567	+4.372608
Steepener	98.699959	97.579381	-1.120578
Flattener	98.699959	99.851766	+1.151807

Currency for face 100. Shape anchors are 2/10/30 years: steepener $-10/0/+10$ bp; flattener $+10/0/-10$ bp. Cash flows and settlement remain fixed. P&L is computed before rounding, so subtracting displayed prices may differ by one last-place unit.

Source: author's calculations using the illustrative offline data supplied with the analytics engine.

Table 8.6: Key-rate DV01 and parallel reconciliation for the 2055 bond. Extra digits expose finite-bump agreement, not economic precision.

Measure	Currency for a 1 bp bump
2-year KR DV01	0.0023690644
5-year KR DV01	0.0072006030
10-year KR DV01	0.0396822358
30-year KR DV01	0.1204744674
Parallel DV01	0.1697264000
Sum of KR DV01	0.1697263707
Sum minus parallel (currency)	$-2.93434113 \times 10^{-8}$

Source: author's calculations using the illustrative offline data supplied with the analytics engine.




Figure 8.2: Key-rate DV01 for the illustrative 2055 bond under continuous-zero discounting and Actual/365 timing. The labelled keys are discrete categories; the vertical axis starts at zero.

Source: author's calculations using the illustrative offline data supplied with the analytics engine.

9 Discussion and limitations

The derivations and numerical checks support the stated calculations. Model risk remains: unsuitable assumptions can produce misleading valuations.

The coupon schedule assumes regular annual or semiannual payments and excludes irregular first or final periods. There is no holiday calendar, business-day adjustment or embedded optionality; accrual and Actual/365 timing are simplified. YTM compresses maturity-dependent discounting into one rate. The assigned-zero interpretation in section 2.1 lacks instrument-based inference, unlike institutional estimation studies [6, 7].

The sub-basis-point RMSE measures agreement with the ten inputs used for fitting, not error on unseen observations or statistical uncertainty. Likewise, noiseless synthetic recovery checks the numerical procedure within its own model family, not that family's adequacy for market instruments. The 30-year key's large sensitivity reflects the distant redemption and late coupons, with their exposure allocated by the chosen basis.

Local calibration may miss better fits; weak identification permits substantially different parameters with similar curves. Neither arbitrage restrictions nor extrapolation validity are established. Deterministic shocks have no estimated probabilities and do not describe stochastic rate dynamics. Invented inputs support no market inference. External production-system validation and market microstructure are absent; trading or backtesting conclusions would exceed the evidence. These limits constrain economic interpretation without invalidating the conditional mathematical identities.

10 Conclusion

The answers to the four research questions depend on the stated assumptions. Explicit coupon and compounding conventions yield coherent prices and analytical sensitivities. Stable NSS loadings and deterministic calibration represent a smooth curve, without ensuring unique parameters. Zero discounting and interpolated shocks give consistent curve exposures, with first-order key-rate additivity. Independent identities and numerical comparisons support the implementation at tested inputs. The study covers valuation, term-structure representation and deterministic risk under these assumptions; it does not establish empirical or production validity.

References

- [1] B. Tuckman and A. Serrat. *Fixed Income Securities: Tools for Today's Markets*. 3rd ed. John Wiley & Sons, 2011.
- [2] U.S. Department of the Treasury. *Interest Rates – Frequently Asked Questions*. URL: <https://home.treasury.gov/policy-issues/financing-the-government/interest-rate-statistics/interest-rates-frequently-asked-questions> (visited on 15/09/2026).
- [3] Bank for International Settlements. *Zero-coupon yield curves: technical documentation*. BIS Papers 25. Bank for International Settlements, 2005. URL: <https://www.bis.org/publ/bppdf/bispap25.htm>.
- [4] C. R. Nelson and A. F. Siegel. ‘Parsimonious Modeling of Yield Curves’. In: *The Journal of Business* 60.4 (1987), pp. 473–489. DOI: 10.1086/296409.
- [5] L. E. O. Svensson. *Estimating and Interpreting Forward Interest Rates: Sweden 1992–1994*. NBER Working Paper 4871. National Bureau of Economic Research, 1994. DOI: 10.3386/w4871. URL: <https://www.nber.org/papers/w4871>.
- [6] N. Anderson and J. Sleath. *New estimates of the UK real and nominal yield curves*. Working Paper 126. Bank of England, 2001. URL: <https://www.bankofengland.co.uk/working-paper/2001/new-estimates-of-the-uk-real-and-nominal-yield-curves>.
- [7] R. S. Gürkaynak, B. Sack and J. H. Wright. ‘The U.S. Treasury yield curve: 1961 to the present’. In: *Journal of Monetary Economics* 54.8 (2007), pp. 2291–2304. DOI: 10.1016/j.jmoneco.2007.06.029.
- [8] F. R. Macaulay. *Some Theoretical Problems Suggested by the Movements of Interest Rates, Bond Yields and Stock Prices in the United States since 1856*. National Bureau of Economic Research, 1938. URL: <https://www.nber.org/books/maca38-1>.
- [9] N. J. Higham. *Accuracy and Stability of Numerical Algorithms*. 2nd ed. Society for Industrial and Applied Mathematics, 2002. DOI: 10.1137/1.9780898718027.
- [10] M. Gilli, S. Große and E. Schumann. *Calibrating the Nelson–Siegel–Svensson model*. COMISEF Working Papers Series WPS-031. Computational Optimization Methods in Statistics, Econometrics and Finance, 30th Mar. 2010. URL: https://www.uni-giessen.de/static_files/pcms/jlu/comisef/files/wps031.pdf.
- [11] T. S. Y. Ho. ‘Key Rate Durations: Measures of Interest Rate Risks’. In: *The Journal of Fixed Income* 2.2 (1992), pp. 29–44. DOI: 10.3905/jfi.1992.408049.

- [12] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith, R. Kern, M. Picus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane, J. Fernández del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Sheppard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke and T. E. Oliphant. ‘Array programming with NumPy’. In: *Nature* 585.7825 (2020), pp. 357–362. DOI: 10.1038/s41586-020-2649-2.
- [13] P. Virtanen, R. Gommers, T. E. Oliphant, M. Haberland, T. Reddy, D. Cournapeau, E. Burovski, P. Peterson, W. Weckesser, J. Bright, S. J. van der Walt, M. Brett, J. Wilson, K. J. Millman, N. Mayorov, A. R. J. Nelson, E. Jones, R. Kern, E. Larson, C. J. Carey, Í. Polat, Y. Feng, E. W. Moore, J. VanderPlas, D. Laxalde, J. Perktold, R. Cimrman, I. Henriksen, E. A. Quintero, C. R. Harris, A. M. Archibald, A. H. Ribeiro, F. Pedregosa, P. van Mulbregt and SciPy 1.0 Contributors. ‘SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python’. In: *Nature Methods* 17.3 (2020), pp. 261–272. DOI: 10.1038/s41592-019-0686-2.
- [14] SciPy developers. *least_squares: SciPy reference documentation*. Online manual inspected: version 1.18.0; executed environment: version 1.18.1. URL: https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.least_squares.html (visited on 14/09/2026).
- [15] M. A. Branch, T. F. Coleman and Y. Li. ‘A Subspace, Interior, and Conjugate Gradient Method for Large-Scale Bound-Constrained Minimization Problems’. In: *SIAM Journal on Scientific Computing* 21.1 (1999), pp. 1–23. DOI: 10.1137/S1064827595289108.
- [16] J. Nocedal and S. J. Wright. *Numerical Optimization*. 2nd ed. Springer, 2006. DOI: 10.1007/978-0-387-40065-5.

A Detailed mathematical derivations

The arguments below restate the independent verification record using the notation introduced at first use in the main chapters. Equalities are over real arithmetic unless a remainder is displayed. Cash flows, settlement and schedules remain fixed when rates change.

A.1 The par-bond identity

On a coupon date, $\alpha = 0$, $q_i = i$ and $AI = 0$. If $c = y_{\text{YTM}}$, then $C = F(B-1)$. Thus equation (3.2) becomes

$$P = F(B-1) \sum_{i=1}^N B^{-i} + FB^{-N}. \quad (\text{A.1})$$

For $B \neq 1$, the geometric sum telescopes:

$$\begin{aligned} (B-1) \sum_{i=1}^N B^{-i} &= \sum_{i=1}^N (B^{1-i} - B^{-i}) \\ &= (1 - B^{-1}) + (B^{-1} - B^{-2}) + \dots + (B^{1-N} - B^{-N}) \\ &= 1 - B^{-N}. \end{aligned}$$

Substitution into equation (A.1) gives $P = F(1 - B^{-N}) + FB^{-N} = F$. For $B = 1$, $c = y_{\text{YTM}} = 0$; coupons vanish and undiscounted redemption gives F . Clean and dirty prices coincide because accrual is zero.

More generally, at coupon-date settlement,

$$P - F = \frac{F(c - y_{\text{YTM}})}{m} \sum_{i=1}^N B^{-i}. \quad (\text{A.2})$$

Consequently coupon above YTM implies premium, and coupon below YTM implies discount, at this settlement boundary. Between coupon dates with $c = y_{\text{YTM}}$, every present value acquires a common multiplier B^α . Hence $P_{\text{dirty}} = FB^\alpha$ and $P_{\text{clean}} = FB^\alpha - C\alpha$, which need not equal F . Fractional compounding and linear accrual are different operations.

A.2 Price derivatives and duration

Because $B = 1 + y_{\text{YTM}}/m > 0$, real powers are differentiable even for noninteger q_i . The chain rule gives

$$\frac{dB^{-q_i}}{dy_{\text{YTM}}} = -q_i B^{-q_i-1} \frac{dB}{dy_{\text{YTM}}} = -\frac{q_i}{m} B^{-q_i-1}.$$

Differentiating the finite sum term by term, using $T_i = q_i/m$ and $PV_i = CF_i B^{-q_i}$,

$$P' = -\sum_i CF_i \frac{q_i}{m} B^{-q_i-1} = -\frac{1}{B} \sum_i T_i PV_i = -P \frac{D_{\text{Mac}}}{B}. \quad (\text{A.3})$$

Dividing by $-P$ proves equation (3.4). The weights PV_i/P are nonnegative and sum to one, so Macaulay duration is a weighted mean of payment times. Positive redemption ensures $P' < 0$. Accrued interest is independent of YTM at fixed settlement; clean price has the same unnormalised derivatives, but normalising by clean price would define different sensitivities.

Differentiating again supplies another $1/m$ and changes the sign:

$$\begin{aligned} P'' &= -\sum_i CF_i \frac{q_i}{m} (-q_i - 1) B^{-q_i-2} \frac{1}{m} \\ &= \sum_i CF_i \frac{q_i(q_i + 1)}{m^2} B^{-q_i-2}. \end{aligned}$$

Normalisation gives the form evaluated by the engine:

$$\mathcal{C} = \sum_i \frac{PV_i}{P} \frac{q_i}{mB} \frac{q_i + 1}{mB} > 0. \quad (\text{A.4})$$

Both factors of $1/m$ express coupon periods in years. The factor $1/2$ belongs to Taylor's second-order price term, not the definition of convexity. Taylor's theorem at y_{YTM} , divided by P , with $P'/P = -D_{\text{Mod}}$ and $P''/P = \mathcal{C}$, yields equation (3.6) with remainder $O((\delta y)^3)$.

A.3 Central differences and DV01

Let $P^{(p)}$ denote the p -th derivative with respect to annual decimal YTM, for positive integer order p . For sufficiently small h , both perturbed yields remain inside $B > 0$. Taylor expansion about the same yield gives

$$P(y_{\text{YTM}} \pm h) = P \pm hP' + \frac{h^2}{2} P'' \pm \frac{h^3}{6} P^{(3)} + \frac{h^4}{24} P^{(4)} + O(h^5).$$

Subtracting the upward price from the downward price cancels the even terms:

$$\begin{aligned} \frac{P(y_{\text{YTM}} - h) - P(y_{\text{YTM}} + h)}{2} &= -hP' - \frac{h^3}{6} P^{(3)} + O(h^5) \\ &= hPD_{\text{Mod}} - \frac{h^3}{6} P^{(3)} + O(h^5). \end{aligned} \quad (\text{A.5})$$

This proves an $O(h^3)$ absolute finite-bump correction to duration-based DV01. Since

$$P^{(3)} = -\sum_i CF_i \frac{q_i(q_i + 1)(q_i + 2)}{m^3} B^{-q_i-3} < 0,$$

the leading correction is positive. Division by two produces a currency response for the stipulated bump; division by $2h$ would instead produce a derivative per unit decimal yield.

The central derivative estimates in tables 7.1 and 7.2 obey

$$-\frac{P(y_{\text{YTM}} + h) - P(y_{\text{YTM}} - h)}{2hP} = D_{\text{Mod}} - \frac{h^2 P^{(3)}}{6P} + O(h^4), \quad (\text{A.6})$$

$$\frac{P(y_{\text{YTM}} + h) - 2P + P(y_{\text{YTM}} - h)}{h^2 P} = \mathcal{C} + \frac{h^2 P^{(4)}}{12P} + O(h^4). \quad (\text{A.7})$$

Each has second-order truncation error. Machine rounding introduces additional error when nearby prices are subtracted; division by a very small h , especially h^2 , amplifies it. Decreasing the step indefinitely need not improve accuracy [9]. The independent report's references are $D_{\text{Mod}} = 4.231517692026542$ and $\mathcal{C} = 21.110444849612590$; the tables preserve that record rather than treating engine output as its own reference.

A.4 NSS limits and stable loadings

Write $x = T/\tau > 0$ for either decay scale. Expanding the exponential,

$$e^{-x} = 1 - x + \frac{x^2}{2} - \frac{x^3}{6} + \frac{x^4}{24} - \frac{x^5}{120} + \frac{x^6}{720} + O(x^7).$$

Subtracting from one, dividing by x , and then subtracting e^{-x} establishes

$$L_1(T; \tau) = 1 - \frac{x}{2} + \frac{x^2}{6} - \frac{x^3}{24} + \frac{x^4}{120} - \frac{x^5}{720} + O(x^6), \quad (\text{A.8})$$

$$L_2(T; \tau) = \frac{x}{2} - \frac{x^2}{3} + \frac{x^3}{8} - \frac{x^4}{30} + \frac{x^5}{144} + O(x^6). \quad (\text{A.9})$$

The constant terms give $L_1 \rightarrow 1$ and $L_2 \rightarrow 0$ as $T \downarrow 0$, proving the short limit in equation (4.3). These are limits, not evaluations at zero: the public curve interface requires positive maturity.

At large T , fixed positive decay scales give $e^{-T/\tau} \rightarrow 0$ and $0 < L_1(T; \tau) < \tau/T \rightarrow 0$. Hence $L_2 \rightarrow 0$ and $r_{\text{NSS}} \rightarrow \beta_0$. More precisely,

$$r_{\text{NSS}}(T) = \beta_0 + \frac{(\beta_1 + \beta_2)\tau_1 + \beta_3\tau_2}{T} + \text{exponentially decaying terms}. \quad (\text{A.10})$$

The short-limit offset β_1 is not the derivative at zero. The linear terms in equations (A.8) and (A.9) give

$$r'_{\text{NSS}}(0^+) = \frac{-\beta_1 + \beta_2}{2\tau_1} + \frac{\beta_3}{2\tau_2},$$

where this prime denotes the maturity derivative, in decimal rate per year, rather than a YTM price derivative.

The small- x polynomials avoid cancellation in both $1 - e^{-x}$ and $L_1 - e^{-x}$. Above the 10^{-3} switch the engine evaluates $L_1 = -\text{expm1}(-x)/x$. The fifth-degree expressions have $O(x^6)$ absolute remainders; the original exponential model remains the mathematical specification.

A.5 Key-rate basis and partition of unity

For ordered keys $0 < k_1 < \dots < k_J$, $J \geq 2$, the endpoint functions are

$$b_1(T) = \begin{cases} 1, & 0 < T \leq k_1, \\ (k_2 - T)/(k_2 - k_1), & k_1 < T < k_2, \\ 0, & T \geq k_2, \end{cases} \quad (\text{A.11})$$

$$b_J(T) = \begin{cases} 0, & T \leq k_{J-1}, \\ (T - k_{J-1})/(k_J - k_{J-1}), & k_{J-1} < T < k_J, \\ 1, & T \geq k_J. \end{cases} \quad (\text{A.12})$$

For an interior key $1 < j < J$,

$$b_j(T) = \begin{cases} (T - k_{j-1})/(k_j - k_{j-1}), & k_{j-1} < T \leq k_j, \\ (k_{j+1} - T)/(k_{j+1} - k_j), & k_j < T < k_{j+1}, \\ 0, & \text{otherwise.} \end{cases} \quad (\text{A.13})$$

At any key only its own weight is one; other weights vanish. Below the first key only $b_1 = 1$, and above the last only $b_J = 1$. Between adjacent keys only their two basis functions are nonzero:

$$b_j(T) + b_{j+1}(T) = \frac{k_{j+1} - T}{k_{j+1} - k_j} + \frac{T - k_j}{k_{j+1} - k_j} = 1. \quad (\text{A.14})$$

These cases exhaust positive maturities and prove equation (5.3). Each weight is nonnegative and continuous, with possibly discontinuous slopes. For one key, $b_1(T) = 1$ everywhere. This proof includes the constant tails, which cannot be inferred from interpolation inside the key range alone.

A.6 First-order and finite-bump reconciliation

Let u_j be a decimal amplitude applied to basis j . With V_i fixed at the base curve,

$$P_{\text{zero}}[r + u_j b_j] = \sum_i V_i e^{-u_j t_i b_j(t_i)}.$$

Differentiating at zero and summing over keys gives

$$\sum_j \left(- \frac{\partial P_{\text{zero}}[r + u_j b_j]}{\partial u_j} \Big|_{u_j=0} \right) = \sum_i V_i t_i \sum_j b_j(t_i) = \sum_i V_i t_i. \quad (\text{A.15})$$

This equals the negative parallel derivative at zero shock in equation (5.2). Differentiating that price again gives $\sum_i t_i^2 V_i e^{-\delta r t_i} > 0$, establishing positive curvature for parallel zero-rate shifts. Multiplication by h proves exact first-order DV01 additivity.

Define bumped curves $r_j^\pm(T) = r(T) \pm h b_j(T)$ and dirty prices $P_j^\pm = P_{\text{zero}}[r_j^\pm]$. For dimensionless argument x , hyperbolic sine is $\sinh(x) = (e^x - e^{-x})/2$. Central repricing gives

$$\text{KRDV01}_j = \frac{P_j^- - P_j^+}{2} = \sum_i V_i \sinh(h t_i b_j(t_i)), \quad (\text{A.16})$$

$$\text{DV01}_\parallel = \sum_i V_i \sinh(h t_i). \quad (\text{A.17})$$

The series $\sinh(x) = x + x^3/6 + x^5/120 + \dots$ contains only odd powers. Substitution into summed key DV01 minus parallel DV01 cancels linear terms by the partition. The resulting cubic term, denoted Δ_3 , is

$$\Delta_{\text{KR}} = \underbrace{\frac{h^3}{6} \sum_i V_i t_i^3 \left(\sum_j b_j(t_i)^3 - 1 \right)}_{\Delta_3} + O(h^5). \quad (\text{A.18})$$

For integer $p > 1$, $0 \leq b_j \leq 1$ implies $b_j^p \leq b_j$, so $\sum_j b_j^p \leq 1$. Every nonzero odd-order correction to the difference is nonpositive. Thus $\Delta_{\text{KR}} \leq 0$ in real arithmetic, strictly negative when at least one positive-value payment is split between neighbouring keys. A single key, or payments supported entirely by individual keys, gives equality.

Define hyperbolic cosine $\cosh(x) = (e^x + e^{-x})/2$. For $x \geq 0$, Taylor's theorem bounds the tail of $\sinh(x)$ after its cubic term by $x^5 \cosh(x)/120$. Each reconciliation coefficient is multiplied by $1 - \sum_j b_j^p \in [0, 1]$; the absolute omitted difference cannot exceed that positive tail. Consequently

$$|\Delta_{\text{KR}} - \Delta_3| \leq \sum_i V_i \frac{(ht_i)^5 \cosh(ht_i)}{120}. \quad (\text{A.19})$$

This bounds the analytic remainder; numerical subtraction adds rounding. For the illustrative risk bond the report records $\Delta_3 = -2.9343412800 \times 10^{-8}$ currency and bound 7.09×10^{-14} . Independent hyperbolic-sine calculations give differences $-2.34747670 \times 10^{-7}$, $-2.93434243 \times 10^{-8}$, and $-3.66792696 \times 10^{-9}$ for $h = 2 \times 10^{-4}, 10^{-4}, 5 \times 10^{-5}$, respectively. Their near-eightfold reductions confirm the cubic leading behaviour. The engine's fixed one-basis-point result in table 8.6 is consistent with this analysis.

A.7 Independent numerical constructions

The derivative reference uses explicit payment dates: 15 July 2025, each January and July in 2026–2029, and 15 January 2030. Settlement is 15 April 2025, giving $\alpha = 90/181$. Nine coupons of 2 precede redemption plus coupon of 102; $m = 2$ and $y_{\text{YTM}} = 0.045$. A separate finite-sum price function, rather than the engine's pricing method, supplies the numerical differences.

For the synthetic NSS check, the known parameter vector is

$$\theta = (0.04, -0.025, 0.03, -0.012, 1.4, 6.5),$$

with coefficient and decay units as in chapter 4. Forty geometrically spaced maturities from 0.05 to 40 years receive rates evaluated from original exponentials using 90-digit decimal arithmetic, then converted to binary64 floating point. Public calibration fits those observations; a further 1,001-point geometric grid checks interpolation. No observation noise is added, although representation rounding remains.

The flat-zero reference uses settlement 1 April 2024, maturity 1 July 2025, coupon 6%, semiannual payments and face 100. Its explicitly specified payments are 3 on 1 July 2024, 3 on 1 January 2025 and 103 at maturity. At continuous annual zero rate 0.045, the independent value is

$$3e^{-0.045(91/365)} + 3e^{-0.045(275/365)} + 103e^{-0.045(456/365)} = 103.23571061465249$$

to displayed precision. No engine-generated schedule enters this reference.

The dense basis check takes the sorted unique union of 100,001 equally spaced maturities from 0.01 to 100 years, keys 2, 5, 10, 30, and midpoints 3.5, 7.5, 20. Independent piecewise formulas avoid calling the implementation's interpolation routine. This distinguishes independent reference construction from merely repeating the same implementation.

B Reproducibility

B.1 Execution record

The dissertation calculations were independently rechecked on 14 September 2026 against the submitted dissertation revision 5a3f2ea. The examiner's findings and supplementary numerical evidence are recorded in `paper/examiner_report.md`. Financial source, tests and illustrative inputs are unchanged from the independent report's audited revision 7eee019. The editable installation imports this repository's package. All 234 tests passed; the individual bond, curve and risk commands completed with the results in chapter 8.

Execution used native Windows PowerShell, Python 3.12.10, NumPy 2.5.3, SciPy 1.18.1, matplotlib 3.11.2, pandas 3.0.5, python-dateutil 2.9.0.post0 and pytest 9.1.1. pandas is declared by the project but is not used in these numerical routines. Dependencies are not locked; deterministic starts ensure same-environment repeatability rather than identical results across all versions.

B.2 Environment and numerical reproduction

Run from the repository root. Create the virtual environment only when absent; installation may need network access, but the examples use only offline inputs.

```
if (-not (Test-Path .\.venv\Scripts\python.exe)) {  
    py -3.12 -m venv .venv  
}  
.\.venv\Scripts\python.exe -m pip install -e ".[dev]"  
.\.venv\Scripts\python.exe -m pytest -q  
.\.venv\Scripts\python.exe -m fixed_income.cli bond  
.\.venv\Scripts\python.exe -m fixed_income.cli curve  
.\.venv\Scripts\python.exe -m fixed_income.cli risk  
.\.venv\Scripts\python.exe -m fixed_income.cli demo
```

Check each exit code before proceeding. The curve and demo commands write `outputs/yield_curve.png`; observations are read from `examples/illustrative_curve.csv` relative to the repository root. Complete independent verification scripts remain in the standalone report and may be executed through its PowerShell stdin wrapper.

B.3 Calibration settings

Initial decay pairs, in years, are (0.5, 2), (1, 5), (2, 10), (5, 1), (10, 3), (3, 15). Each seeds linear least-squares coefficients, followed by joint six-parameter optimisation. Coefficients are unbounded; both decay scales lie in [0.01, 100] years. At least six distinct positive maturities are required; observations have equal weight, including repeated tenors.

Optimisation residuals are $10^4 e_\ell$. The installed `least_squares` defaults are trust-region reflective, linear loss and a two-point numerical Jacobian (a matrix of residual derivatives). The code specifies Jacobian-based scaling, `ftol=xtol=gtol=1e-10` and `max_nfev=3000`. These tolerances control termination through objective, step and gradient criteria; they are not bounds on financial error. Function counts exclude extra residual calls for numerical Jacobians [14].

Only successful, finite, bounded candidates qualify; smallest unscaled residual sum of squares wins, with earlier-start tie breaking. Failure of all starts raises an error. This procedure is a local search, not a guarantee of global optimality.

B.4 Figure and numerical provenance

Figure 8.1 uses the existing fitted curve evaluated at 500 evenly spaced maturities between the minimum and maximum supplied observations, with rates multiplied by 100 for percentage display. Figure 8.2 plots public key-rate DV01 results at the four default keys. Both are vector PDF figures generated using existing NumPy and matplotlib dependencies through stdin Python. The examiner corrected the key-rate figure's clipped axis label; its regeneration script is retained in `paper/examiner_report.md`. No financial module or permanent plotting feature was added. Table values were compared with fresh command outputs or unrounded public-API results; reconciliation was computed before rounding.

B.5 LaTeX compilation

The build uses pdfLaTeX, latexmk 4.88 and Biber 2.22 with compressed numeric biblatex references; pdfTeX is 1.40.29 (MiKTeX 26.5).

```
Push-Location .\paper
try {
    latexmk -C -outdir=build main.tex
    if ($LASTEXITCODE -ne 0) { throw "LaTeX cleanup failed." }
    latexmk -pdf -interaction=nonstopmode -halt-on-error -outdir=build main.tex
    if ($LASTEXITCODE -ne 0) { throw "LaTeX compilation failed." }
    Copy-Item .\build\main.pdf .\fixed_income_analytics_dissertation.pdf
} finally {
    Pop-Location
}
git diff --check
```

Auxiliary files remain in ignored `paper/build/`. The final PDF and two figure PDFs are intentional

publication artefacts.

B.6 Bibliographic and editorial verification

All bibliography entries have verified metadata recorded in `paper/bibliography_verification.md` and are cited in the dissertation. Primary papers support the NSS attribution; institutional studies supply estimation context, with no external financial observations imported. Mathematical identities are derived rather than inferred from bibliographic authority.